A frozen 3D U-Net is adapted to a sequence of brain-MRI tasks from 32 labeled volumes each, with one LoRA module and output head per task. Whether small adapters suffice depends on where the task needs capacity, and the two task types answer oppositely. Segmentation inherits a large decoder, so full-decoder linear probing is marginally better (Dice 0.789 vs 0.777) at 90x the trainable parameters, and the two are equal by 64 shots. Brain age has no such head: given whole-head input, LoRA is the only frozen method that learns it (MAE 10.2 yr, R2 = 0.40) against a 14.5 yr constant baseline, beating full fine-tuning of 14M parameters by 2 years with 56K, while linear probing and BitFit emit a constant. Leaving the 66-parameter output projection frozen costs 0.22 Dice. All results use a subject-grouped held-out protocol with whole-volume inference over three seeds.